

Reference Material

Deng Cai (蔡登)

College of Computer Science
Zhejiang University

dengcai@gmail.com





Three Topics

- ▶ Bayesian Decision Theory
- ▶ Artificial Neural Network & Deep Learning
- ▶ k Nearest Neighbor Classifier

▶ Experimental materials download link:

<http://summer2019.in.zjulearning.org/>

Bayesian Decision Theory

Deng Cai (蔡登)

College of Computer Science
Zhejiang University

dengcai@gmail.com






Goal

- ▶ To finish the problems *bayes_decision_rule* and *text_classification*, you need to know:
 - What is Bayes' Theorem.
 - How to calculate the prior, likelihood, and posterior.
 - What are maximum likelihood decision rule, optimal bayes decision rule (maximum posterior), and minimum bayes risk rule.
How to design a classifier according to the above rules.

Bayesian Decision Theory

- ▶ Decision problem posed in probabilistic terms
- ▶ x : sample
- ▶ ω : state of the nature
- ▶ $P(\omega|x)$: given x , what is the probability of the state of the nature.
- ▶ Sea bass / Salmon Example





Basics of Probability

- ▶ An experiment is a well-defined process with observable outcomes.
- ▶ The set or collection of all outcomes of an experiment is called the sample space, S .
- ▶ An event E is any subset of outcomes from S .
- ▶ Probability of an event, $P(E)$ is $P(E) = \frac{\text{number of outcomes in } E}{\text{number of outcomes in } S}$.



Bayes' Theorem

- ▶ Conditional probability: $P(A|B) = \frac{P(A, B)}{P(B)}$.
 - Test of Independence: A and B are said to be independent if and only if $P(A, B) = P(A) P(B)$.

- ▶ Bayes' Theorem:
$$\text{posterior } P(A|B) = \frac{\text{likelihood } P(B|A) \text{ prior } P(A)}{P(B)}$$



Prior

- ▶ A priori (prior) probability of the state of nature
 - Random variable (State of nature is unpredictable)
 - Reflects our prior knowledge about how likely we are to observe a sea bass or salmon
 - The catch of salmon and sea bass is equiprobable
 - $P(\omega_1) = P(\omega_2)$ (uniform priors)
 - $P(\omega_1) + P(\omega_2) = 1$ (exclusivity and exhaustivity)
- ▶ Decision rule with only the prior information
 - Decide ω_1 if $P(\omega_1) > P(\omega_2)$, otherwise decide ω_2



Likelihood

- ▶ Suppose now we have a measurement or feature on the state of nature - say the fish lightness value
- ▶ $P(x|\omega_1)$ and $P(x|\omega_2)$ describe the difference in lightness feature between populations of sea bass and salmon
- ▶ $P(x|\omega_j)$ is called the **likelihood** of ω_j with respect to x ; the category ω_j for which $P(x | \omega_j)$ is large is more likely to be the true category
- ▶ **Maximum likelihood decision**
 - Assign input pattern x to class ω_1 if
$$P(x | \omega_1) > P(x | \omega_2), \text{ otherwise } \omega_2$$



Posterior

- Bayes formula

$$P(\omega_i|x) = \frac{P(x|\omega_i)P(\omega_i)}{P(x)}$$

$$P(x) = \sum_{i=1}^k P(x|\omega_i)P(\omega_i)$$

- **Posterior** = (**Likelihood** \times **Prior**) / Evidence
 - Evidence $P(x)$ can be viewed as a scale factor that guarantees that the posterior probabilities sum to 1

Posterior \propto Likelihood \times Prior



Optimal Bayes Decision Rule

- ▶ $P(\omega_1 | x)$ is the probability of the state of nature being ω_1 given that feature value x has been observed
- ▶ Decision given the posterior probabilities, **Optimal Bayes Decision rule**

X is an observation for which:

if $P(\omega_1 | x) > P(\omega_2 | x)$ \rightarrow True state of nature = ω_1

if $P(\omega_1 | x) < P(\omega_2 | x)$ \rightarrow True state of nature = ω_2

Bayes decision rule minimizes the probability of error, that is the term **Optimal** comes from. But why? Can you prove it?



Optimal Bayes Decision Rule

Based on Bayes decision rule, whenever we observe a particular x , the probability of error is:

$$P(\text{error} \mid x) = P(\omega_1 \mid x) \text{ if we decide } \omega_2$$

$$P(\text{error} \mid x) = P(\omega_2 \mid x) \text{ if we decide } \omega_1$$

Bayes decision rule:

Decide ω_1 if $P(\omega_1 \mid x) > P(\omega_2 \mid x)$; otherwise decide ω_2

Therefore:

$$P(\text{error} \mid x) = \min [P(\omega_1 \mid x), P(\omega_2 \mid x)]$$

- The unconditional error, $P(\text{error})$, obtained by integration over all x w.r.t. $p(x)$



Optimal Bayes Decision Rule

- ▶ Decide ω_1 if $P(\omega_1 | x) > P(\omega_2 | x)$;
otherwise decide ω_2

- ▶ Special cases:
 - (i) $P(\omega_1) = P(\omega_2)$; Decide ω_1 if
 $P(x | \omega_1) > P(x | \omega_2)$, otherwise ω_2

Maximum likelihood decision

 - (ii) $P(x | \omega_1) = P(x | \omega_2)$; Decide ω_1 if
 $P(\omega_1) > P(\omega_2)$, otherwise ω_2



Bayes Risk

- ▶ Conditional risk

$$R(\alpha_i|\mathbf{x}) = \sum_{j=1}^c \lambda(\alpha_i|\omega_j)P(\omega_j|\mathbf{x})$$

- ▶ Select the action for which the conditional risk $R(\alpha_i|\mathbf{x})$ is *minimum*

$$R = \int R(\alpha_i|\mathbf{x}) p(\mathbf{x}) d\mathbf{x}$$

- ▶ Risk R is minimum and R in this case is called the
 - Bayes risk = best performance that can be achieved!



Example 1: Two-category classification

α_1 : deciding ω_1

α_2 : deciding ω_2

$$\lambda_{ij} = \lambda(\alpha_i \mid \omega_j)$$

Conditional risk:

$$R(\alpha_1 \mid x) = \lambda_{11}P(\omega_1 \mid x) + \lambda_{12}P(\omega_2 \mid x)$$

$$R(\alpha_2 \mid x) = \lambda_{21}P(\omega_1 \mid x) + \lambda_{22}P(\omega_2 \mid x)$$

How to achieve Bayes risk?



Example 1: Two-category classification

Bayes rule is the following:

$$\text{if } R(\alpha_1 \mid x) < R(\alpha_2 \mid x)$$

action α_1 : “decide ω_1 ” is taken

This results in the equivalent rule:

decide ω_1 if:

$$(\lambda_{21} - \lambda_{11}) P(x \mid \omega_1) P(\omega_1) > (\lambda_{12} - \lambda_{22}) P(x \mid \omega_2) P(\omega_2)$$

and decide ω_2 otherwise



Example 1: Two-category classification

- ▶ The preceding rule is equivalent to the following rule:

- ▶ If $\frac{P(\mathbf{x}|\omega_1)}{P(\mathbf{x}|\omega_2)} > \frac{\lambda_{12}-\lambda_{22}}{\lambda_{21}-\lambda_{11}} \times \frac{P(\omega_2)}{P(\omega_1)}$

Then take action α_1 (decide ω_1)

Otherwise take action α_2 (decide ω_2)

- ▶ “If the **likelihood ratio** exceeds a threshold value that is independent of the input pattern \mathbf{x} , we can take optimal actions”



Naïve Bayes Classifier

- ▶ Given $\mathbf{x} = (x_1, \dots, x_p)^T$
 - Goal is to predict class ω
 - Specifically, we want to find the value of ω that maximizes $P(\omega|\mathbf{x}) = P(\omega|x_1, \dots, x_p)$

$$P(\omega|x_1, \dots, x_p) \propto P(x_1, \dots, x_p|\omega)P(\omega)$$

- ▶ Independence assumption among features

$$P(x_1, \dots, x_p|\omega) = P(x_1|\omega) \cdots P(x_p|\omega)$$



How to Estimate Probabilities from Data?

Tid	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

- ▶ Class: $P(\omega_k) = \frac{N_{\omega_k}}{N}$
 - e.g., $P(\text{No}) = 7/10$,
 $P(\text{Yes}) = 3/10$

- ▶ For discrete attributes:

$$P(x_i|\omega_k) = \frac{|x_{ik}|}{N_{\omega_k}}$$

- where $|x_{ik}|$ is number of instances having attribute x_i and belongs to class ω_k
- Examples:

$$P(\text{Status}=\text{Married}|\text{No}) = 4/7$$
$$P(\text{Refund}=\text{Yes}|\text{Yes})=0$$



How to Estimate Probabilities from Data?

- ▶ For continuous attributes:
 - **Discretize** the range into bins
 - one ordinal attribute per bin
 - violates independence assumption
 - **Two-way split:** $(x < v)$ or $(x > v)$
 - choose only one of the two splits as new attribute
 - **Probability density estimation:**
 - Assume attribute follows a normal distribution
 - Use data to estimate parameters of distribution (e.g., mean and standard deviation)
 - Once probability distribution is known, can use it to estimate the conditional probability $P(x_1|\omega)$

How to Estimate Probabilities from Data?

Tid	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

► Normal distribution:

$$P(x_i | \omega_j) = \frac{1}{\sqrt{2\pi\sigma_{ij}^2}} \exp\left(-\frac{(x_i - \mu_{ij})^2}{2\sigma_{ij}^2}\right)$$

■ One for each (x_i, ω_i) pair

► For (Income, Class=No):

■ If Class=No

- sample mean = 110
- sample variance = 2975

$$P(\text{Income} = 120 | \text{No}) = \frac{1}{\sqrt{2\pi}(54.54)} \exp\left(-\frac{(120 - 110)^2}{2(2975)}\right) = 0.0072$$



Example of Naïve Bayes Classifier

Given a Test Record:

$X = (\text{Refund} = \text{No}, \text{Married}, \text{Income} = 120\text{K})$

naive Bayes Classifier:

$$P(\text{Refund}=\text{Yes}|\text{No}) = 3/7$$

$$P(\text{Refund}=\text{No}|\text{No}) = 4/7$$

$$P(\text{Refund}=\text{Yes}|\text{Yes}) = 0$$

$$P(\text{Refund}=\text{No}|\text{Yes}) = 1$$

$$P(\text{Marital Status}=\text{Single}|\text{No}) = 2/7$$

$$P(\text{Marital Status}=\text{Divorced}|\text{No}) = 1/7$$

$$P(\text{Marital Status}=\text{Married}|\text{No}) = 4/7$$

$$P(\text{Marital Status}=\text{Single}|\text{Yes}) = 2/7$$

$$P(\text{Marital Status}=\text{Divorced}|\text{Yes}) = 1/7$$

$$P(\text{Marital Status}=\text{Married}|\text{Yes}) = 0$$

For taxable income:

If class=No: sample mean=110
 sample variance=2975

If class=Yes: sample mean=90
 sample variance=25

- $P(X|\text{Class}=\text{No}) = P(\text{Refund}=\text{No}|\text{Class}=\text{No})$
 $\times P(\text{Married}|\text{Class}=\text{No})$
 $\times P(\text{Income}=120\text{K}|\text{Class}=\text{No})$
 $= 4/7 \times 4/7 \times 0.0072 = 0.0024$
- $P(X|\text{Class}=\text{Yes}) = P(\text{Refund}=\text{No}|\text{Class}=\text{Yes})$
 $\times P(\text{Married}|\text{Class}=\text{Yes})$
 $\times P(\text{Income}=120\text{K}|\text{Class}=\text{Yes})$
 $= 1 \times 0 \times 1.2 \times 10^{-9} = 0$

Since $P(X|\text{No})P(\text{No}) > P(X|\text{Yes})P(\text{Yes})$

Therefore $P(\text{No}|X) > P(\text{Yes}|X)$

$\Rightarrow \text{Class} = \text{No}$

Example of Naïve Bayes Classifier

Name	Give Birth	Can Fly	Live in Water	Have Legs	Class
human	yes	no	no	yes	mammals
python	no	no	no	no	non-mammals
salmon	no	no	yes	no	non-mammals
whale	yes	no	yes	no	mammals
frog	no	no	sometimes	yes	non-mammals
komodo	no	no	no	yes	non-mammals
bat	yes	yes	no	yes	mammals
pigeon	no	yes	no	yes	non-mammals
cat	yes	no	no	yes	mammals
leopard shark	yes	no	yes	no	non-mammals
turtle	no	no	sometimes	yes	non-mammals
penguin	no	no	sometimes	yes	non-mammals
porcupine	yes	no	no	yes	mammals
eel	no	no	yes	no	non-mammals
salamander	no	no	sometimes	yes	non-mammals
gila monster	no	no	no	yes	non-mammals
platypus	no	no	no	yes	mammals
owl	no	yes	no	yes	non-mammals
dolphin	yes	no	yes	no	mammals
eagle	no	yes	no	yes	non-mammals
Give Birth	Can Fly	Live in Water	Have Legs	Class	
yes	no	yes	no	?	

A: attributes

M: mammals

N: non-mammals

$$P(A \mid M) = \frac{6}{7} \times \frac{6}{7} \times \frac{2}{7} \times \frac{2}{7} = 0.06$$

$$P(A \mid N) = \frac{1}{13} \times \frac{10}{13} \times \frac{3}{13} \times \frac{4}{13} = 0.0042$$

$$P(A \mid M)P(M) = 0.06 \times \frac{7}{20} = 0.021$$

$$P(A \mid N)P(N) = 0.004 \times \frac{13}{20} = 0.0027$$

$$P(A|M)P(M) > P(A|N)P(N)$$

=> Mammals



Naïve Bayes (Summary)

► Advantages

- Robust to isolated noise points
- Handle missing values by ignoring the instance during probability estimate calculations
- Robust to irrelevant attributes

► Disadvantages

- Independence assumption may not hold for some attributes
- Smoothing

$$P(x_i|\omega_k) = \frac{|x_{ik}| + 1}{N_{\omega_k} + K}$$

Artificial Neural Network & Deep Learning

Deng Cai (蔡登)

College of Computer Science
Zhejiang University

dengcai@gmail.com





Goal


- ▶ To finish the problem *neural_networks*, you need to know:
 - What is artificial neural network.
 - The forward and backpropagation processes of the neural network.




Natural Neural Net Models

- ▶ Human brain consists of very large number of neurons (between 10^{10} to 10^{12})
- ▶ No. of interconnections per neuron is between 1K to 10K
- ▶ Total number of interconnections is about 10^{14}
- ▶ Damage to a few neurons or synapse (links) does not appear to impair overall performance significantly (robustness)


The Artificial Neural Network



- ▶ A cartoon drawing of a biological neuron



The Artificial Neural Network



- This is the fully-connected neural network. There are many varieties of neural network, *e.g.*, convolutional neural network (CNN) and recurrent neural network (RNN).



Activation Function

- ▶ **Activation Function f**
 - Must be non-linear (otherwise, 3-layer network is just a linear discriminant) and saturate (have max and min value) to keep weights and activation functions bounded
 - Activation function and its derivative must be continuous and smooth; optionally monotonic
 - Choice may depend on the problem. Eg. Gaussian activation if the data comes from a mixture of Gaussians
 - Eg: sigmoid (most popular), polynomial, tanh

- ▶ **Parameters of activation function (e.g. Sigmoid)**
 - Centered at 0, odd function $f(-\text{net}) = -f(\text{net})$ (anti-symmetric); leads to faster learning
 - Depend on the range of the input values



Activation Function

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parametric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

.....



General Feedforward Operation

- ▶ Case of c output units

$$g_k(\mathbf{x}) \equiv z_k = f \left(\sum_{j=1}^{n_H} w_{kj} f \left(\sum_{i=1}^d w_{ji} x_i + w_{j0} \right) + w_{k0} \right)$$

$$k = 1, \dots, c$$

- ▶ Hidden units enable us to express more complicated nonlinear functions and extend classification capability
- ▶ Assume for now that all activation functions are identical
- ▶ Question: Can every decision boundary be implemented by a three-layer network described by the above equation?



Expressive Power of Multi-layer Networks

- ▶ Answer: Yes (due to A. Kolmogorov)
 - Any continuous function from input to output can be implemented in a three-layer net, **given sufficient number of hidden units n_H , proper nonlinearities, and weights.**
- ▶ Any continuous function $g(\mathbf{x})$ defined on the unit hypercube I^n ($I = [0,1]$ and $n \geq 2$) can be represented in the following form:

$$g(\mathbf{x}) = \sum_{j=1}^{2n+1} \Xi_j \left(\sum_{i=1}^d \Phi_{ij}(x_i) \right)$$


for properly chosen functions Ξ_j and Φ_{ij}



Gradient Descent



Backpropagation Algorithm




$$\Delta \mathbf{w} = -\eta \nabla J = -\eta \frac{\partial J}{\partial \mathbf{w}} \quad \Delta w_{mn} = -\eta \frac{\partial J}{\partial w_{mn}}$$

- Where η is the learning rate which indicates the relative size of the change in weights

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \Delta \mathbf{w}^{(t)}$$

- where t indexes the particular pattern presentation

Backpropagation Algorithm




$$J = \frac{1}{2} \sum_{k=1}^c (t_k - z_k)^2$$

$$\frac{\partial J}{\partial w_{kj}} = \frac{\partial J}{\partial z_k} \cdot \frac{\partial z_k}{\partial net_k} \cdot \frac{\partial net_k}{\partial w_{kj}} = (z_k - t_k) \cdot f'(net_k) \cdot y_j$$

$$\begin{aligned} \frac{\partial J}{\partial w_{ji}} &= \frac{\partial J}{\partial y_j} \cdot \frac{\partial y_j}{\partial net_j} \cdot \frac{\partial net_j}{\partial w_{ji}} \\ &= \left(\sum_{k=1}^c (z_k - t_k) \cdot f'(net_k) \cdot w_{kj} \right) \cdot f'(net_j) \cdot x_i \end{aligned}$$

Backpropagation Algorithm




vanishing gradient problem

$$\frac{\partial J}{\partial w_4} = \frac{\partial J}{\partial z} \cdot f'(net_4) \cdot x_4$$

$$\frac{\partial J}{\partial w_3} = \frac{\partial J}{\partial z} \cdot f'(net_4) \cdot w_4 \cdot f'(net_3) \cdot x_3$$

$$\frac{\partial J}{\partial w_2} = \frac{\partial J}{\partial z} \cdot f'(net_4) \cdot w_4 \cdot f'(net_3) \cdot w_3 f'(net_2) \cdot x_2$$

$$\frac{\partial J}{\partial w_1} = \frac{\partial J}{\partial z} \cdot \overset{< 1}{f'(net_4) \cdot w_4} \cdot \overset{< 1}{f'(net_3) \cdot w_3} \cdot \overset{< 1}{f'(net_2) \cdot w_2} \cdot \overset{< 1}{f'(net_1)} \cdot x_1$$



First order derivative

k Nearest Neighbor Classifier

Deng Cai (蔡登)

College of Computer Science
Zhejiang University

dengcai@gmail.com



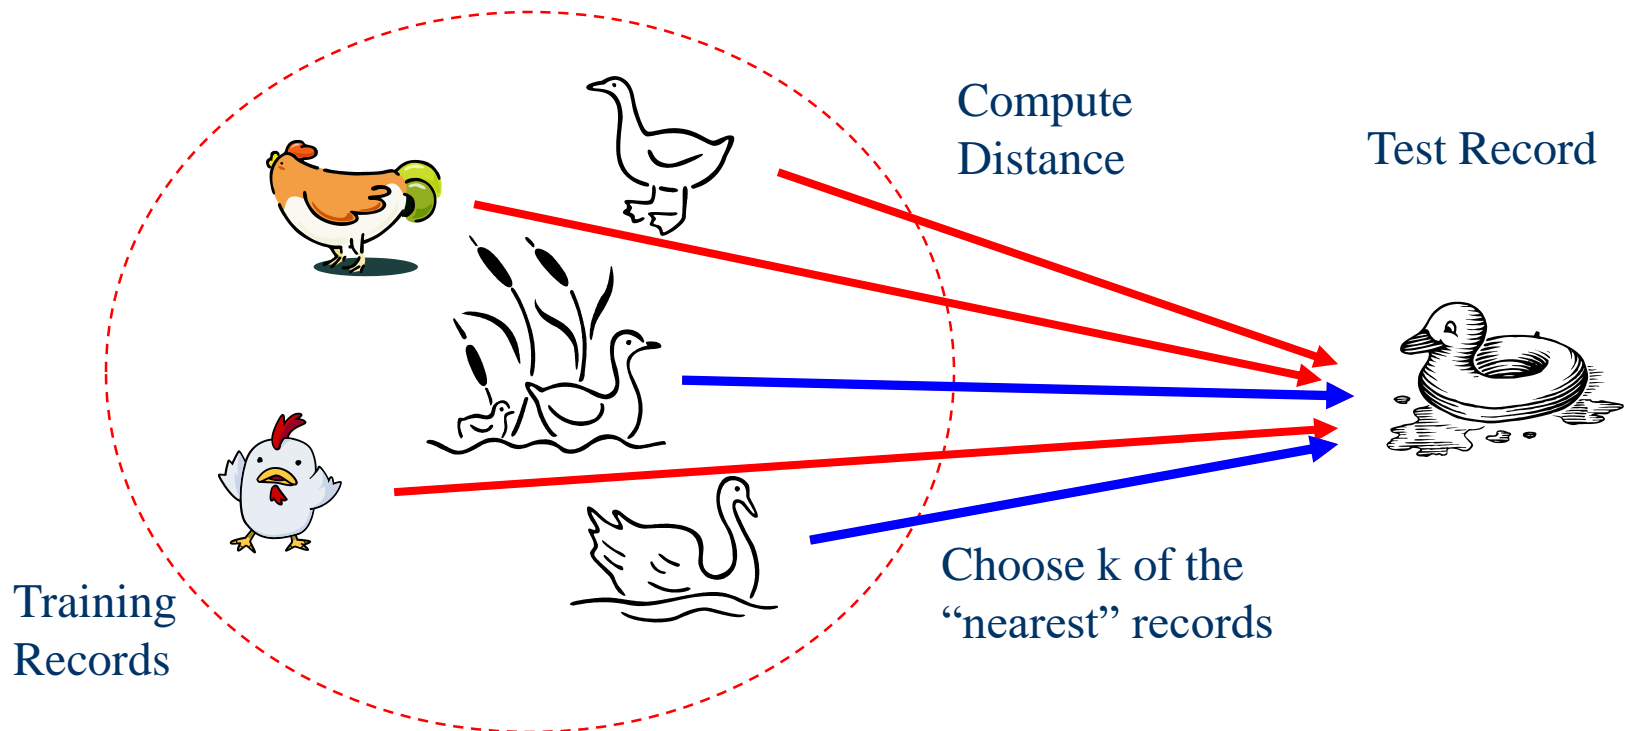


Goal

- ▶ To finish the problem *knn*, you need to know:
 - What is the main idea of *knn*.
 - How to apply *knn* in a practical problem.


Nearest Neighbor Classifiers

- ▶ Basic idea:
 - If it walks like a duck, quacks like a duck, then it's probably a duck




Nearest-Neighbor Classifiers

Unknown record




- Requires three things
 - The set of stored records
 - Distance Metric to compute distance between records
 - The value of k , the number of nearest neighbors to retrieve
- To classify an unknown record:
 - Compute distance to other training records
 - Identify k nearest neighbors
 - Use class labels of nearest neighbors to determine the class label of unknown record (e.g., by taking majority vote)


Definition of Nearest Neighbor



(a) 1-nearest neighbor



(b) 2-nearest neighbor



(c) 3-nearest neighbor

K -nearest neighbors of a record x are data points that have the k smallest distance to x



How many parameters in kNN?

- ▶ A Linear Classifier

$$f(x) = w^T x$$


- The number of parameters?

- ▶ kNN Classifier

- **Effective** number of parameters?

$$\frac{N}{k}$$

15-Nearest Neighbor Classifier



1-Nearest Neighbor Classifier

